


```

grunt> dump student;
2020-11-25 04:26:36.978 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2020-11-25 04:26:37.136 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 04:26:37.553 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 04:26:37.648 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 04:26:37.648 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 04:26:37.962 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:26:38.733 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 04:26:38.929 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.reduce.markreset.buffer.percent is deprecated. Instead, use mapreduce.reduce.markreset.buffer.percent
2020-11-25 04:26:38.929 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 04:26:38.929 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.output.compress is deprecated. Instead, use mapreduce.output.fileoutputformat.compress
2020-11-25 04:26:56.356 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:1
2020-11-25 04:26:57.414 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301019383_0006
2020-11-25 04:26:58.439 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301019383_0006
2020-11-25 04:26:58.783 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8080/proxy/application_1606301019383_0006/
2020-11-25 04:26:58.783 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301019383_0006
2020-11-25 04:26:58.783 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases student
2020-11-25 04:26:58.784 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: student[1,10],student[-1,-1]
2020-11-25 04:26:58.784 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:56010/jobdetail.jsp?jobid=job_1606301019383_0006
2020-11-25 04:26:58.997 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 04:27:52.581 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 04:27:54.733 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.reduce.tasks is deprecated. Instead, use mapreduce.job.reduces
2020-11-25 04:27:54.916 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 04:27:54.923 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 04:26:38 2020-11-25 04:27:54 UNKNOWN

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
ature Outputs
job_1606301019383_0006 1 0 19 19 19 19 n/a n/a n/a n/a student MAP_ONLY hdfs://quickstart.cloudera:8028/tmp/temp-1340583213/tmp553182092,

Input(s):
Successfully read 9 records (589 bytes) from: "/user/cloudera/trainP/students.txt"

Output(s):
Successfully stored 9 records (216 bytes) in: "hdfs://quickstart.cloudera:8028/tmp/temp-1340583213/tmp553182092"

Counters:
Total records written : 9
Total bytes written : 216
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301019383_0006

2020-11-25 04:27:55.460 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 04:27:55.473 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:27:55.473 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:27:55.479 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 04:27:55.531 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:27:55.531 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(1,sushil,68.2,F)
(2,deepak,70.33,F)
(3,siddhesh,80.21,F)
(4,vikas,78.3,F)
(5,sumit,62.0,F)
(6,ankita,68.0,F)
(7,supriya,72.21,F)
(8,neha,91.22,F)
(9,nayuresh,85.66,F)
grunt> █

```

5. Projection:-

```

grunt> studentname = foreach student generate name;
grunt> dump studentname;
2020-11-25 04:31:16.512 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2020-11-25 04:31:16.515 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 04:31:16.531 [main] INFO org.apache.pig.newplan.logical.rules.ColumnPruneVisitor - Columns pruned for student: $0, $2, $3
2020-11-25 04:31:16.555 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 04:31:16.558 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 04:31:16.558 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 04:31:16.645 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:31:16.652 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 04:31:16.682 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3

```

```

2020-11-25 04:31:17,551 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job42684880511787951.jar
2020-11-25 04:31:18,427 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job42684880511787951.jar created
2020-11-25 04:31:30,534 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 04:31:30,536 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 04:31:30,536 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 04:31:30,537 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 04:31:30,609 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 04:31:30,609 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:31:30,625 [JobControl] INFO org.apache.hadoop.yarn.client.WMProxy - Connecting to ResourceManager at /%0.0.0.0:8032
2020-11-25 04:31:30,644 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:31:31,258 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:31:31,259 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 04:31:31,265 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 04:31:31,363 [DataStreamer] INFO org.apache.hadoop.mapreduce.lib.output.FileOutputFormat - Total output paths to process : 1
2020-11-25 04:31:31,363 [DataStreamer] WARN org.apache.hadoop.hdfs.DFSClient - Caught exception
java.lang.InterruptedException
    at java.lang.Object.wait(Native Method)
    at java.lang.Thread.join(Thread.java:1281)
    at java.lang.Thread.join(Thread.java:1355)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.closeResponder(DFSOutputStream.java:952)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.endBlock(DFSOutputStream.java:690)
    at org.apache.hadoop.hdfs.DFSOutputStreamDataStreamer.run(DFSOutputStream.java:879)
2020-11-25 04:31:31,363 [JobControl] INFO org.apache.hadoop.mapreduce.lib.output.FileOutputFormat - Total output paths to process : 1
2020-11-25 04:31:31,687 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0007
2020-11-25 04:31:31,916 [JobControl] INFO org.apache.hadoop.yarn.client.impl.YarnClientImpl - Submitted application application_1606301819383_0007
2020-11-25 04:31:31,938 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8088/proxy/application_1606301819383_0007/
2020-11-25 04:31:31,938 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0007
2020-11-25 04:31:31,931 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases student, studentname
2020-11-25 04:31:31,931 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: student[1,30], studentname[1,1] C: R:
2020-11-25 04:31:31,931 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
ails.jsp?jobid=job_1606301819383_0007
2020-11-25 04:31:32,177 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 04:32:20,627 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 04:32:22,689 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 04:32:22,638 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 04:31:30 2020-11-25 04:32:22 UNKNOWNV

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
ature Outputs
job_1606301819383_0007 1 0 19 19 19 19 n/a n/a n/a n/a student,studentname MAP_ONLY hdfs://quickstart.cloude
ra:8020/tmp/temp-1348583213/tmp1948339100

Input(s):
Successfully read 9 records (589 bytes) from: "/user/cloudera/trainP/students.txt"

Output(s):
Successfully stored 9 records (118 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1348583213/tmp1948339100"

Counters:
Total records written : 9
Total bytes written : 118
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0007

2020-11-25 04:32:22,853 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 04:32:22,854 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:32:22,854 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:32:22,860 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set., will not generate code.
2020-11-25 04:32:22,878 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:32:22,878 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
{sushil}
{deepak}
{siddhesh}
{vikas}
{sumit}
{ankita}
{supriya}
{neha}
{mayuresh}
grint> █

```

6. Joins-

```

grunt> cat /user/cloudera/trainP/dept.txt
10,sushil,engineer,50000
10,nikhil,engineer,45000
10,supriya,engineer,50000
10,siddesh,engineer,44000
20,manish,clerk,20000
20,mahesh,clerk,25000
30,minal,scientist,60000
30,krishna,scientist,80000
30,govind,scientist,60000
50,rahul,chemist,80000
40,neha,biochemist,90000
grunt> █

```



```

grunt> dept = load '/user/cloudera/train@dept.txt' USING PigStorage(), ('asid:int,name:chararray,deptname:chararray,sal:int');
grunt> dump dept;
2020-11-25 04:37:31,874 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2020-11-25 04:37:31,876 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 04:37:31,897 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MRCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 04:37:31,910 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 04:37:31,910 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 04:37:32,037 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:37:32,041 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 04:37:32,058 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 04:37:32,893 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - creating jar file Job1918742779818018191.jar
2020-11-25 04:37:44,679 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - jar file Job1918742779818018191.jar created
2020-11-25 04:37:44,700 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - Setting up single store job
2020-11-25 04:37:44,701 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 04:37:44,701 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 04:37:44,701 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 04:37:44,742 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 04:37:44,742 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:37:44,746 [JobControl] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:37:44,796 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:37:45,587 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:37:45,588 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths to process : 1
2020-11-25 04:37:45,592 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths (combined) to process : 1
2020-11-25 04:37:45,664 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:1
2020-11-25 04:37:45,830 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0010
2020-11-25 04:37:45,937 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0010
2020-11-25 04:37:45,943 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8088/proxy/application_1606301819383_0010/
2020-11-25 04:37:45,943 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - HadoopJobId: job_1606301819383_0010
2020-11-25 04:37:45,943 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - Processing aliases dept
2020-11-25 04:37:45,943 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - detailed locations: M: dept[5,7],dept[-1,-1] C: R
2020-11-25 04:37:45,943 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
ails.jsp?jobid=job_1606301819383_0010
2020-11-25 04:37:46,042 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - 0% complete
2020-11-25 04:38:37,083 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - 50% complete
2020-11-25 04:38:41,429 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - 100% complete
2020-11-25 04:38:41,429 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion PigVersion UserID StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 04:37:32 2020-11-25 04:38:41 UNKNOWN

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
ature Outputs
job_1606301819383_0010 1 8 21 21 21 21 n/a n/a n/a n/a dept MAP_ONLY hdfs://quickstart.cloudera:8020/tmp/temp
-1340503213/tmp906094038,

Input(s):
Successfully read 11 records (652 bytes) from: "/user/cloudera/train@dept.txt"

Output(s):
Successfully stored 11 records (334 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp906094038"

Counters:
Total records written : 11
Total bytes written : 334
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0010

2020-11-25 04:38:41,671 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MapReduceLauncher - Success!
2020-11-25 04:38:41,672 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:38:41,672 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:38:41,673 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 04:38:41,688 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:38:41,688 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths to process : 1
(10,sushil,engineer,50000)
(10,nikhil,engineer,45000)
(10,supriya,engineer,50000)
(10,siddesh,engineer,44000)
(20,manish,clerk,20000)
(20,mahesh,clerk,25000)
(30,minal,scientist,60000)
(30,krishna,scientist,80000)
(30,govind,scientist,60000)
(50,rahul,chemist,80000)
(40,neha,biochemist,90000)
grunt> █

```

7. Relational Operators –

- Cross:** The cross operator is used to calculate the cross product of two or more relations.

```

grunt> x = cross student,dept;
grunt> dump x;
2020-11-25 04:42:07,381 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: CROSS
2020-11-25 04:42:07,382 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 04:42:07,434 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MRCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 04:42:07,472 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MRCompilerLastInputStreamingOptimizer - Rewrite: POPackage->POForEach to POJoinPackage
2020-11-25 04:42:07,474 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 04:42:07,474 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 04:42:07,566 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:42:07,571 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 04:42:07,611 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 04:42:07,612 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - Reduce phase detected, estimating # of required reducers.
2020-11-25 04:42:07,615 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.InputSizeReducerEstimator
2020-11-25 04:42:07,624 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.InputSizeReducerEstimator - BytesPerReducer:1000000000 maxReducers:999
totalInputFileSize=477
2020-11-25 04:42:07,624 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduce_layer.JobControlCompiler - Setting Parallelism to 1

```

```

2020-11-25 04:42:08,636 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job25487612432153817.jar
2020-11-25 04:42:22,829 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job25487612432153817.jar created
2020-11-25 04:42:22,851 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 04:42:22,854 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 04:42:22,854 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 04:42:22,854 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting Key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 04:42:22,276 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 04:42:22,276 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:42:22,276 [JobControl] INFO org.apache.hadoop.yarn.client.NMProxy - Connecting to ResourceManager at /s0.0.0:8032
2020-11-25 04:42:22,357 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:42:22,728 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:42:22,730 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 04:42:22,730 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 04:42:22,741 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:42:22,741 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 04:42:22,745 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 04:42:22,829 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:2
2020-11-25 04:42:22,996 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0011
2020-11-25 04:42:23,149 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0011
2020-11-25 04:42:23,155 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8088/proxy/application_1606301819383_0011/
2020-11-25 04:42:23,155 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0011
2020-11-25 04:42:23,155 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases dept,student,x
2020-11-25 04:42:23,156 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Detailed locations: M: dept[5,7],dept[-1,-1],x[6,4],x[6,4],student[1,10],student[-1,-1],x[6,4],x[6,4] C: R:
2020-11-25 04:42:23,156 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50830/jobdetalls.jsp?jobid=job_1606301819383_0011
2020-11-25 04:42:23,272 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 04:43:42,068 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 04:44:14,029 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 04:44:14,838 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 04:42:07 2020-11-25 04:44:14 CR055

Success!

Job Stats (time in seconds):
JobId Maps Reduces MapMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1606301819383_0011 2 1 43 43 43 43 25 25 25 25 dept,student,x hdfs://quickstart.cloudera:8020/tmp/temp-1340803213/tmp1606060015,

Input(s):
Successfully read 11 records from: "/user/cloudera/trainP/dept.txt"
Successfully read 9 records from: "/user/cloudera/trainP/students.txt"

Output(s):
Successfully stored 99 records (4906 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340803213/tmp1606060015"
Counters:
Total records written : 99
Total bytes written : 4906
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0011

2020-11-25 04:44:15,095 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 04:44:15,096 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:44:15,096 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:44:15,098 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 04:44:15,114 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 04:44:15,114 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(9,nayuresh,85.66,F,40,neha,biochemist,90000)
(9,nayuresh,85.66,F,50,rahul,chemist,80000)
(9,nayuresh,85.66,F,30,govind,scientist,60000)
(9,nayuresh,85.66,F,30,krishna,scientist,80000)
(9,nayuresh,85.66,F,30,minal,scientist,60000)
(9,nayuresh,85.66,F,20,mahesh,clerk,25000)
(9,nayuresh,85.66,F,20,manish,clerk,20000)
(9,nayuresh,85.66,F,10,siddesh,engineer,44000)
(9,nayuresh,85.66,F,10,supriya,engineer,50000)
(9,nayuresh,85.66,F,10,nikhil,engineer,45000)
(9,nayuresh,85.66,F,10,sushil,engineer,50000)
(8,neha,91.22,F,40,neha,biochemist,90000)
(8,neha,91.22,F,50,rahul,chemist,80000)
(8,neha,91.22,F,30,govind,scientist,60000)
(8,neha,91.22,F,30,krishna,scientist,80000)
(8,neha,91.22,F,30,minal,scientist,60000)
(8,neha,91.22,F,20,mahesh,clerk,25000)
(8,neha,91.22,F,20,manish,clerk,20000)
(8,neha,91.22,F,10,siddesh,engineer,44000)
(8,neha,91.22,F,10,supriya,engineer,50000)
(8,neha,91.22,F,10,nikhil,engineer,45000)
(8,neha,91.22,F,10,sushil,engineer,50000)
(7,supriya,72.21,F,40,neha,biochemist,90000)
(7,supriya,72.21,F,50,rahul,chemist,80000)
(7,supriya,72.21,F,30,govind,scientist,60000)
(7,supriya,72.21,F,30,krishna,scientist,80000)
(7,supriya,72.21,F,30,minal,scientist,60000)
(7,supriya,72.21,F,20,mahesh,clerk,25000)
(7,supriya,72.21,F,20,manish,clerk,20000)
(7,supriya,72.21,F,10,siddesh,engineer,44000)
(7,supriya,72.21,F,10,supriya,engineer,50000)
(7,supriya,72.21,F,10,nikhil,engineer,45000)
(7,supriya,72.21,F,10,sushil,engineer,50000)
(6,ankita,68.0,F,40,neha,biochemist,90000)
(6,ankita,68.0,F,50,rahul,chemist,80000)
(6,ankita,68.0,F,30,govind,scientist,60000)
(6,ankita,68.0,F,30,krishna,scientist,80000)
(6,ankita,68.0,F,30,minal,scientist,60000)
(6,ankita,68.0,F,20,mahesh,clerk,25000)
(6,ankita,68.0,F,20,manish,clerk,20000)
(6,ankita,68.0,F,10,siddesh,engineer,44000)
(6,ankita,68.0,F,10,supriya,engineer,50000)
(6,ankita,68.0,F,10,nikhil,engineer,45000)
(6,ankita,68.0,F,10,sushil,engineer,50000)
(5,sumit,62.0,F,40,neha,biochemist,90000)
(5,sumit,62.0,F,50,rahul,chemist,80000)
(5,sumit,62.0,F,30,govind,scientist,60000)
(5,sumit,62.0,F,30,krishna,scientist,80000)
(5,sumit,62.0,F,30,minal,scientist,60000)
(5,sumit,62.0,F,20,mahesh,clerk,25000)
(5,sumit,62.0,F,20,manish,clerk,20000)
(5,sumit,62.0,F,10,siddesh,engineer,44000)
(5,sumit,62.0,F,10,supriya,engineer,50000)

```



```
(5,sumit,62,0,F,10,nikhil,engineer,45000)
(5,sumit,62,0,F,10,sushil,engineer,50000)
(4,vikas,70,3,F,40,neha,biochemist,90000)
(4,vikas,70,3,F,50,rahu,chemist,80000)
(4,vikas,70,3,F,30,govind,scientist,60000)
(4,vikas,70,3,F,30,krishna,scientist,80000)
(4,vikas,70,3,F,30,minal,scientist,60000)
(4,vikas,70,3,F,20,mahesh,clerk,25000)
(4,vikas,70,3,F,20,manish,clerk,20000)
(4,vikas,70,3,F,10,siddesh,engineer,44000)
(4,vikas,70,3,F,10,supriya,engineer,50000)
(4,vikas,70,3,F,10,nikhil,engineer,45000)
(4,vikas,70,3,F,10,sushil,engineer,50000)
(3,siddesh,80,21,F,40,neha,biochemist,90000)
(3,siddesh,80,21,F,50,rahu,chemist,80000)
(3,siddesh,80,21,F,30,govind,scientist,60000)
(3,siddesh,80,21,F,30,krishna,scientist,80000)
(3,siddesh,80,21,F,30,minal,scientist,60000)
(3,siddesh,80,21,F,20,mahesh,clerk,25000)
(3,siddesh,80,21,F,20,manish,clerk,20000)
(3,siddesh,80,21,F,10,siddesh,engineer,44000)
(3,siddesh,80,21,F,10,supriya,engineer,50000)
(3,siddesh,80,21,F,10,nikhil,engineer,45000)
(3,siddesh,80,21,F,10,sushil,engineer,50000)
(2,deepak,70,33,F,40,neha,biochemist,90000)
(2,deepak,70,33,F,50,rahu,chemist,80000)
(2,deepak,70,33,F,30,govind,scientist,60000)
(2,deepak,70,33,F,30,krishna,scientist,80000)
(2,deepak,70,33,F,30,minal,scientist,60000)
(2,deepak,70,33,F,20,mahesh,clerk,25000)
(2,deepak,70,33,F,20,manish,clerk,20000)
(2,deepak,70,33,F,10,siddesh,engineer,44000)
(2,deepak,70,33,F,10,supriya,engineer,50000)
(2,deepak,70,33,F,10,nikhil,engineer,45000)
(2,deepak,70,33,F,10,sushil,engineer,50000)
(1,sushil,80,2,F,40,neha,biochemist,90000)
(1,sushil,80,2,F,50,rahu,chemist,80000)
(1,sushil,80,2,F,30,govind,scientist,60000)
(1,sushil,80,2,F,30,krishna,scientist,80000)
(1,sushil,80,2,F,30,minal,scientist,60000)
(1,sushil,80,2,F,20,mahesh,clerk,25000)
(1,sushil,80,2,F,20,manish,clerk,20000)
(1,sushil,80,2,F,10,siddesh,engineer,44000)
(1,sushil,80,2,F,10,supriya,engineer,50000)
(1,sushil,80,2,F,10,nikhil,engineer,45000)
(1,sushil,80,2,F,10,sushil,engineer,50000)
grunt>
```

- b. **Distinct**: This operator is used to remove the duplicate tuples in a relation. It does not preserve the original order of the contents.

```
grunt> cat /user/cloudera/trainP/dummyf.txt
1,2,2,3,4
2,4,5,65
2,4,5,65
23,44,3,2
23,44,3,2
1,2,3,4
4,3,2,1
5,6,7,8
grunt>
```

```
grunt> A = load '/user/cloudera/trainP/dummyf.txt' USING PigStorage(',') as (a1:int,a2:int,a3:int,a4:int);
2020-11-25 04:52:04,620 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:52:04,621 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
grunt> dump A;
2020-11-25 04:52:16,221 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2020-11-25 04:52:16,222 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushDownFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 04:52:16,242 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MRCCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 04:52:16,245 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 04:52:16,245 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 04:52:16,335 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:52:16,339 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 04:52:16,354 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 04:52:19,379 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.JobControlCompiler - jar file Job526854400995090202.jar created
2020-11-25 04:52:19,406 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 04:52:19,401 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 04:52:19,401 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting process to move generated code to distributed cache
2020-11-25 04:52:19,401 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 04:52:19,437 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 04:52:19,437 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:52:19,441 [JobControl] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:52:19,476 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:52:19,824 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:52:19,825 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 04:52:19,828 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 04:52:19,902 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Number of splits:1
2020-11-25 04:52:30,243 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0012
2020-11-25 04:52:30,379 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0012
2020-11-25 04:52:30,387 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8088/proxy/application_1606301819383_0012/
2020-11-25 04:52:30,387 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0012
2020-11-25 04:52:30,387 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - Processing aliases A
2020-11-25 04:52:30,387 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - Detailed locations: M: A[7,4],A[1,1] C: R:
2020-11-25 04:52:30,388 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - More information at http://localhost:50030/jobdet
ails.jsp?jobid=job_1606301819383_0012
2020-11-25 04:52:30,541 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - 0% complete
2020-11-25 04:52:30,548 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - 50% complete
2020-11-25 04:52:30,548 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreduceLayer.MapReduceLauncher - 100% complete
2020-11-25 04:52:30,548 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:
```

```

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 04:52:16 2020-11-25 04:53:26 UNKNOWN

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
easure Outputs
job 1606301819383 0012 1 0 22 22 22 22 n/a n/a n/a n/a A MAP_ONLY hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp151831464,

Input(s):
Successfully read 8 records (454 bytes) from: "/user/cloudera/trainP/dummyf.txt"

Output(s):
Successfully stored 8 records (93 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp151831464"

Counters:
Total records written : 8
Total bytes written : 93
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job 1606301819383 0012
2020-11-25 04:53:26.464 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 04:53:26.465 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:53:26.465 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.job.tracker.address
2020-11-25 04:53:26.466 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set,... will not generate code.
2020-11-25 04:53:26.482 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:53:26.482 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(1,2,3)
(2,4,5,65)
(2,4,5,65)
(23,44,3,2)
(23,44,3,2)
(1,2,3,4)
(4,3,2,1)
(5,6,7,8)
grunt> █

grunt> z = distinct A;
grunt> dump z;
2020-11-25 04:57:03.162 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: DISTINCT
2020-11-25 04:57:03.162 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColu
mRewrite, GroupByCostParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForE
achFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionsSimplifier, PartitionFilterOptimizer]}
2020-11-25 04:57:03.185 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 04:57:03.214 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 04:57:03.214 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 04:57:03.280 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 04:57:03.290 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 04:57:03.307 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not
set, set to default 0.3
2020-11-25 04:57:03.388 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required r
educers.
2020-11-25 04:57:03.388 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.h
adoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2020-11-25 04:57:03.312 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer=1000000000 MaxReducers=999
TotalInputFileSize=72
2020-11-25 04:57:03.312 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1
2020-11-25 04:57:17.906 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 04:57:17.910 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 04:57:17.910 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 04:57:17.910 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 04:57:17.912 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting identity combiner class.
2020-11-25 04:57:17.953 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 04:57:17.954 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.job.tracker.address
2020-11-25 04:57:18.000 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:57:18.346 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:57:18.346 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 04:57:18.350 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 04:57:18.530 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job 1606301819383 0013
2020-11-25 04:57:18.626 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application 1606301819383 0013
2020-11-25 04:57:18.632 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8088/proxy/application_1606301819383_0
013/
2020-11-25 04:57:18.632 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job 1606301819383 0013
2020-11-25 04:57:18.632 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases A
2020-11-25 04:57:18.632 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: A[7,4],A[-1,-1] C: R:
2020-11-25 04:57:18.632 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
ails.jsp?jobid=job 1606301819383 0013
2020-11-25 04:57:18.711 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 04:58:07.244 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 04:58:35.117 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 04:58:35.117 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 04:57:03 2020-11-25 04:58:35 DISTINCT

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
easure Outputs
job 1606301819383 0013 1 1 20 20 20 20 21 21 21 21 A DISTINCT hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp-680608746,

Input(s):
Successfully read 8 records (454 bytes) from: "/user/cloudera/trainP/dummyf.txt"

Output(s):
Successfully stored 8 records (69 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp-680608746"

Counters:
Total records written : 8
Total bytes written : 69
Spillable Memory Manager spill count : 0

```



```
Total bags proactively spilled: 0
Total records proactively spilled: 0
```

```
Job DAG:
job 1606301819383_0014
```

```
2020-11-25 04:58:35.368 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 04:58:35.369 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 04:58:35.370 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 04:58:35.371 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 04:58:35.383 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 04:58:35.383 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(1,2,2,3)
(1,2,3,4)
(2,4,5,65)
(4,3,2,1)
(5,6,7,8)
(23,44,3,2)
grunt>
```

c. **Filter:** This operator can be used to select the required data based on some conditions

```
grunt> y = FILTER A By a2 == 2;
2020-11-25 05:02:15.030 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:02:15.031 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
grunt> dump y;
2020-11-25 05:02:20.877 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: FILTER
2020-11-25 05:02:20.878 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED:[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED:[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 05:02:20.880 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 05:02:20.883 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 05:02:20.887 [main] INFO org.apache.hadoop.yarn.client.NMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:02:20.892 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 05:02:20.236 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 05:02:20.976 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job443133628658263935.jar
2020-11-25 05:02:32.784 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job443133628658263935.jar created
2020-11-25 05:02:32.805 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 05:02:32.806 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 05:02:32.806 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 05:02:32.832 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 05:02:32.882 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:02:32.892 [JobControl] INFO org.apache.hadoop.yarn.client.NMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:02:32.929 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:02:33.235 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:02:33.235 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:02:33.239 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 05:02:33.453 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0014
2020-11-25 05:02:33.594 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0014
2020-11-25 05:02:33.682 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8080/proxy/application_1606301819383_0014/
2020-11-25 05:02:33.693 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0014
2020-11-25 05:02:33.693 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases A,y
2020-11-25 05:02:33.693 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: A[7,4],A[-1,-1],y[9,4] C: K:
2020-11-25 05:02:33.693 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:5083/jobdet
ails.jsp?jobid=job_1606301819383_0014
2020-11-25 05:02:33.758 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 05:02:33.470 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 05:03:29.808 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to j
ob history server
2020-11-25 05:03:31.399 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 05:03:31.406 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:
```

```
HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 05:02:28 2020-11-25 05:03:31 FILTER

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1606301819383_0014 1 0 20 20 20 20 n/a n/a n/a n/a A,y MAP_ONLY hdfs://quickstart.cloudera:8020/tmp/temp-1340583213/tp1340908926

Input(s):
Successfully read 0 records (054 bytes) from: "/user/cloudera/trainP/dummy.txt"

Output(s):
Successfully stored 2 records (22 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340583213/tp1340908926"

Counters:
Total records written : 2
Total bytes written : 22
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0014
```

```
2020-11-25 05:03:31.623 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 05:03:31.624 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:03:31.624 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:03:31.625 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 05:03:31.640 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:03:31.648 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(1,2,2,3)
(1,2,3,4)
grunt>
```

Grunt>y = FILTER dept by deptname == 'engineer';
Grunt> dump y;


```

grunt> y = FILTER dept by deptname == 'engineer')
grunt> dump y;
2020-11-25 05:06:38,207 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: FILTER
2020-11-25 05:06:38,208 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - (RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInsert, LimitOptimizer, LoadTypeCastInsert, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInsert], RULES_DISABLED=[FilterLogicExpressionsSimplifier, PartitionFilterOptimizer])
2020-11-25 05:06:38,226 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MRCCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 05:06:38,230 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 05:06:38,230 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 05:06:38,551 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:06:38,572 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 05:06:39,836 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - creating jar file Job6338823264492229631.jar
2020-11-25 05:06:51,940 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - jar file Job6338823264492229631.jar created
2020-11-25 05:06:51,969 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - Setting up single store job
2020-11-25 05:06:51,978 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code
2020-11-25 05:06:51,978 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 05:06:51,978 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 05:06:52,825 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 05:06:52,825 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:06:52,830 [JobControl] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:06:52,859 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:06:52,789 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:06:52,789 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:06:52,792 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 05:06:52,867 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:1
2020-11-25 05:06:52,960 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0015
2020-11-25 05:06:53,888 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0015
2020-11-25 05:06:53,890 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8080/proxy/application_1606301819383_0015/
2020-11-25 05:06:53,892 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0015
2020-11-25 05:06:53,892 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - Processing aliases dept,y
2020-11-25 05:06:53,892 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - detailed locations: M: dept[5,7],dept[-1,-1],y[11,4],C: #;
2020-11-25 05:06:53,892 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
all: [jobId=job_1606301819383_0015]
2020-11-25 05:06:53,192 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - 0% complete
2020-11-25 05:07:43,090 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - 50% complete
2020-11-25 05:07:40,656 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - 100% complete
2020-11-25 05:07:49,057 [main] INFO org.apache.pig.tools.pigstats.SamplePigStats - Script Statistics:

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 05:06:38 2020-11-25 05:07:49 FILTER

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1606301819383_0015 1 0 21 21 21 21 n/a n/a n/a n/a dept,y MAP_ONLY hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp-1562494189,

Input(s):
Successfully read 11 records (652 bytes) from: "/user/cloudera/trainP/dept.txt"

Output(s):
Successfully stored 4 records (126 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp-1562494189"

Counters:
Total records written : 4
Total bytes written : 126
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0015

2020-11-25 05:07:49,289 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - Success!
2020-11-25 05:07:49,293 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:07:49,293 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:07:49,293 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 05:07:49,315 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:07:49,315 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
(18,sushil,engineer,50000)
(18,nikhil,engineer,49000)
(18,supriya,engineer,50000)
(18,siddesh,engineer,44000)
grunt> █

```

d. **FOREACH**:- This operator is used to generate data transformation based on column data

```

grunt> x = foreach dept GENERATE id,name;
grunt> dump x;
2020-11-25 05:09:54,073 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2020-11-25 05:09:54,075 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - (RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInsert, LimitOptimizer, LoadTypeCastInsert, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushUpFilter, SplitFilter, StreamTypeCastInsert], RULES_DISABLED=[FilterLogicExpressionsSimplifier, PartitionFilterOptimizer])
2020-11-25 05:09:54,078 [main] INFO org.apache.pig.newplan.logical.rules.ColumnPruneVisitor - Columns pruned for dept: $2, $3
2020-11-25 05:09:54,092 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MRCCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 05:09:54,095 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 05:09:54,095 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 05:09:54,175 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:09:54,179 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 05:09:54,198 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 05:09:55,726 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - creating jar file Job1724687739560467346.jar
2020-11-25 05:10:07,677 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - jar file Job1724687739560467346.jar created
2020-11-25 05:10:07,698 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.JobControlCompiler - Setting up single store job
2020-11-25 05:10:07,699 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 05:10:07,699 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 05:10:07,699 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 05:10:07,733 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapreducelayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 05:10:07,733 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:10:07,736 [JobControl] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:10:07,746 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:10:08,032 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:10:08,032 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:10:08,036 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1

```

```

2020-11-25 05:10:08.108 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:1
2020-11-25 05:10:08.273 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0016
2020-11-25 05:10:08.419 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0016
2020-11-25 05:10:08.427 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8088/proxy/application_1606301819383_0016/
2020-11-25 05:10:08.427 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0016
2020-11-25 05:10:08.428 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases X,dept
2020-11-25 05:10:08.428 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: dept[S,Y]X[-1,-1] C: R:
2020-11-25 05:10:08.428 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
ails.jsp?jobid=job_1606301819383_0016
2020-11-25 05:10:08.574 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 05:10:36.510 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 05:11:04.829 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 05:11:04.830 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 05:09:54 2020-11-25 05:11:04 UNKNOWN

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1606301819383_0016 1 0 21 21 21 21 n/a n/a n/a n/a X,dept MAP_ONLY hdfs://quickstart.cloudera:8020/tmp/tmp-
-1340583213/tmp-2007464689,

Input(s):
Successfully read 11 records (652 bytes) from: "/user/cloudera/trainP/dept.txt"

Output(s):
Successfully stored 11 records (164 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/tmp-1340583213/tmp-2007464689"

Counters:
Total records written : 11
Total bytes written : 164
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0016

2020-11-25 05:11:04.262 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 05:11:04.263 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:11:04.263 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:11:04.265 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 05:11:04.279 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:11:04.279 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
{10,sushill}
{10,nikhil}
{10,supriya}
{10,siddesh}
{20,manish}
{20,mahesh}
{30,minal}
{30,krishna}
{30,govind}
{30,rahul}
{40,naha}
grust=

```

e. COGROUP operator:-

```

grunt> copyFromLocal /home/cloudera/trainingPracs/dummyfb.txt /user/cloudera/trainP
grunt> cat /user/cloudera/trainP/dummyfb.txt

```

```

1,2
2,3
3,4
5,6
65,3
3,6
4,3
23,44
grunt>

```

```

grunt> B = load '/user/training/dumB.txt' USING PigStorage(',') as (b1:int, b2:int);
grunt> dump B;

```

```

grunt> B = load '/user/cloudera/trainP/dummyfb.txt' USING PigStorage(',') as (b1:int,b2:int);
grunt> dump B;
2020-11-25 05:14:38.287 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: UNKNOWN
2020-11-25 05:14:38.289 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColu
mnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForE
achFlatten, PushupFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 05:14:38.380 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 05:14:38.385 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 05:14:38.385 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 05:14:38.378 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:14:38.381 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 05:14:38.388 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not
set, set to default 0.3

```



```

2020-11-25 05:14:39,197 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job2842863976494419855.jar
2020-11-25 05:14:39,351 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job2842863976494419855.jar created
2020-11-25 05:14:39,377 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 05:14:39,381 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 05:14:39,381 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 05:14:39,382 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 05:14:39,419 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 05:14:39,426 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:14:39,442 [JobControl] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:14:39,493 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:14:39,587 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:14:39,587 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths to process : 1
2020-11-25 05:14:39,585 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths (combined) to process : 1
2020-11-25 05:14:39,590 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:1
2020-11-25 05:14:39,593 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0018
2020-11-25 05:14:39,603 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0018
2020-11-25 05:14:39,608 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8080/proxy/application_1606301819383_0018/
2020-11-25 05:14:39,609 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0018
2020-11-25 05:14:39,609 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases B
2020-11-25 05:14:39,609 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: B[14,4],R[-1,-1] C=: R:
2020-11-25 05:14:39,609 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50038/jobdet
ails.jsp?jobid=job_1606301819383_0018
2020-11-25 05:14:39,716 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 05:15:43,202 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 30% complete
2020-11-25 05:15:50,077 [main] INFO org.apache.hadoop.mapred.ClientServiceDelegate - Application state is completed. FinalApplicationStatus=SUCCEEDED. Redirecting to j
ob history server
2020-11-25 05:15:50,927 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 05:15:50,927 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:
HadoopVersion PigVersion UserId StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 05:14:38 2020-11-25 05:15:50 UNKNOWN

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
ature Outputs
job_1606301819383_0018 1 0 20 20 20 20 n/a n/a n/a n/a B MAP ONLY hdfs://quickstart.cloudera:8020/tmp/temp-1340583213/tmp-1062959807,
-1340583213/tmp-1062959807,

Input(s):
Successfully read 0 records (410 bytes) from: "/user/cloudera/train@dummyfb.txt"

Output(s):
Successfully stored 0 records (03 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340583213/tmp-1062959807"

Counters:
Total records written : 0
Total bytes written : 03
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0018
2020-11-25 05:15:51,515 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 05:15:51,516 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:15:51,516 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:15:51,542 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 05:15:51,589 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:15:51,589 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapReduceUtil - Total input paths to process : 1
(1,2)
(2,3)
(3,4)
(5,6)
(65,3)
(3,6)
(4,3)
(23,44)
grunt> █

```

Grunt> C = COGROUP A By a1 inner, B BY b1 Inner;

Grunt>Dump C;

```

grunt> C = COGROUP A By a1 inner, B BY b1 Inner;
grunt> dump C;
2020-11-25 05:17:40,078 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: COGROUP
2020-11-25 05:17:40,088 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColu
mnRewrite, GroupByKeyParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushdownForE
achFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 05:17:40,119 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 100 optimistic? False
2020-11-25 05:17:40,127 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 05:17:40,127 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 05:17:40,211 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:17:40,215 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 05:17:40,271 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not
set, set to default 0.3
2020-11-25 05:17:40,272 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required r
educers.
2020-11-25 05:17:40,275 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.h
adoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2020-11-25 05:17:40,295 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer=1000000000 maxReducers=999
totalInputFileSize=107
2020-11-25 05:17:40,295 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1

```



```

2020-11-25 05:17:42,055 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job3174296207502596419.jar
2020-11-25 05:17:54,300 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job3174296207502596419.jar created
2020-11-25 05:17:54,328 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 05:17:54,330 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 05:17:54,338 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 05:17:54,338 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 05:17:54,386 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 05:17:54,386 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:17:54,398 [JobControl] INFO org.apache.hadoop.yarn.client.NMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:17:54,422 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:17:55,143 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:17:55,144 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:17:55,179 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:17:55,179 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:17:55,188 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 05:17:55,273 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0019
2020-11-25 05:17:55,517 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1606301819383_0019
2020-11-25 05:17:55,523 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8088/proxy/application_1606301819383_0019/
2020-11-25 05:17:55,523 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0019
2020-11-25 05:17:55,523 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases A,B,C
2020-11-25 05:17:55,523 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: A[7,4],A[-1,-1],C[10,4],0[1,4],0[-1,-1],C[15,4] C: R:
2020-11-25 05:17:55,523 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
ails.jsp?jobid=job_1606301819383_0019
2020-11-25 05:17:55,618 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 05:19:04,218 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 05:19:31,761 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 05:19:31,762 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:
HadoopVersion PigVersion Userid StartedAt FinishedAt Features
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 05:17:40 2020-11-25 05:19:31 COGROUP

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1606301819383_0019 2 1 40 40 40 40 21 21 21 21 A,B,C COGROUP hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp-117546071,
213/tmp-117546071.

Input(s):
Successfully read 8 records from: "/user/cloudera/train/p/dummyfs.txt"
Successfully read 8 records from: "/user/cloudera/train/p/dummyf.txt"

Output(s):
Successfully stored 5 records (142 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp-117546071"

Counters:
Total records written : 5
Total bytes written : 142
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0019
2020-11-25 05:19:32,249 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 05:19:32,254 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:19:32,254 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:19:32,260 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 05:19:32,290 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:19:32,290 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
{(1,(1,2,3,4)),(1,(2,2,3))},{(1,3)}
{(2,(2,4,5,6)),(2,4,5,6)},{(2,3)}
{(4,(4,3,2,1))},{(4,3)}
{(5,(5,6,7,8))},{(5,6)}
{(23,((23,44,3,2),(23,44,3,2))),(23,44)}
grunt> █

```

8. Nested Projection:

If one of the field in input relation is a [tuple,bag] then we can perform projection operation on that field.

Grunt>Z=FOREACH C GENERATE group, B.b2;

Grunt>dump Z;

```

grunt> Z = foreach C GENERATE group, B.b2;
grunt> dump Z;
2020-11-25 05:21:28,828 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: COGROUP
2020-11-25 05:21:28,832 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColu
mRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForE
achFlatten, PushUpFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionSimplifier, PartitionFilterOptimizer]}
2020-11-25 05:21:28,840 [main] INFO org.apache.pig.newplan.logical.rules.ColumnPruneVisitor - columns pruned for A: $1, $2, $3
2020-11-25 05:21:28,865 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCompiler - File concatenation threshold: 1000 optimistic? false
2020-11-25 05:21:28,870 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size before optimization: 1
2020-11-25 05:21:28,870 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - MR plan size after optimization: 1
2020-11-25 05:21:28,961 [main] INFO org.apache.hadoop.yarn.client.NMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:21:28,965 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 05:21:29,000 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not
set, set to default 0.3
2020-11-25 05:21:29,007 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required r
educers.
2020-11-25 05:21:29,007 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.h
adoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2020-11-25 05:21:29,015 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer=1000000000 maxReducers=999
totalInputFileSize=167
2020-11-25 05:21:29,015 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1

```

```

2020-11-25 05:21:38,653 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job7071070210231782646.jar
2020-11-25 05:21:42,503 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job7071070210231782646.jar created
2020-11-25 05:21:42,523 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 05:21:42,525 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 05:21:42,525 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 05:21:42,525 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 05:21:42,568 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:21:42,574 [JobControl] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:21:42,593 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:21:43,258 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:21:43,258 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 05:21:43,269 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:21:43,270 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:21:43,279 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 05:21:43,351 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:2
2020-11-25 05:21:43,543 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1006301819383_0020
2020-11-25 05:21:43,682 [JobControl] INFO org.apache.hadoop.yarn.client.api.impl.YarnClientImpl - Submitted application application_1006301819383_0020
2020-11-25 05:21:43,687 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8080/proxy/application_1006301819383_0020/
2020-11-25 05:21:43,687 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1006301819383_0020
2020-11-25 05:21:43,687 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases A,B,C,Z
2020-11-25 05:21:43,687 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: A[7,4],A[-1,-1],C[15,4],B[14,4],B[-1,-1],C[15,4] C: B: 2[-1,-1]
2020-11-25 05:21:43,687 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
all.jsp?jobid=job_1006301819383_0020
2020-11-25 05:21:43,870 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 05:22:51,843 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 05:23:20,272 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 05:23:20,273 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion: PigVersion: UserId: StartedAt: FinishedAt: Features:
2.6.0-cdh5.12.0 0.12.0-cdh5.12.0 cloudera 2020-11-25 05:21:38 2020-11-25 05:23:20 COGROUP

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
job_1006301819383_0020 2 1 38 38 38 38 21 21 21 21 A,B,C,Z COGROUP hdfs://quickstart.cloudera:8020/tmp/temp-1340583213/tmp1784299764,

Input(s):
Successfully read 8 records from: "/user/cloudera/train@dummyf.txt"
Successfully read 8 records from: "/user/cloudera/train@dummyfb.txt"

Output(s):
Successfully stored 5 records (54 bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340583213/tmp1784299764"

Counters:
Total records written : 5
Total bytes written : 54
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1006301819383_0020

2020-11-25 05:23:20,524 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 05:23:20,525 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:23:20,525 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:23:20,527 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 05:23:20,541 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:23:20,541 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
{1,{12}}
{2,{13}}
{4,{13}}
{5,{16}}
{23,{44}}
grunt>

```

Grunt>Z = FOREACH C GENERATE group,A.(a1,a2);
Grunt>dump Z;

```

grunt> Z = foreach C GENERATE group,A.(a1,a2)
grunt> dump Z;
2020-11-25 05:26:22,884 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig features used in the script: COGROUP
2020-11-25 05:26:22,888 [main] INFO org.apache.pig.newplan.logical.optimizer.LogicalPlanOptimizer - {RULES_ENABLED=[AddForEach, ColumnMapKeyPrune, DuplicateForEachColumnRewrite, GroupByConstParallelSetter, ImplicitSplitInserter, LimitOptimizer, LoadTypeCastInserter, MergeFilter, MergeForEach, NewPartitionFilterOptimizer, PushDownForEachFlatten, PushupFilter, SplitFilter, StreamTypeCastInserter], RULES_DISABLED=[FilterLogicExpressionsSimplifier, PartitionFilterOptimizer]}
2020-11-25 05:26:22,110 [main] INFO org.apache.pig.newplan.logical.rules.ColumnPruneVisitor - Columns pruned for B: $1
2020-11-25 05:26:22,166 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MRCCompiler - File concatenation threshold: 100 optimistic? false
2020-11-25 05:26:22,177 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - RM plan size before optimization: 1
2020-11-25 05:26:22,177 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MultiQueryOptimizer - RM plan size after optimization: 1
2020-11-25 05:26:22,354 [main] INFO org.apache.hadoop.yarn.client.RMProxy - Connecting to ResourceManager at /0.0.0.0:8032
2020-11-25 05:26:22,361 [main] INFO org.apache.pig.tools.pigstats.ScriptState - Pig script settings are added to the job
2020-11-25 05:26:22,443 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - mapred.job.reduce.markreset.buffer.percent is not set, set to default 0.3
2020-11-25 05:26:22,445 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Reduce phase detected, estimating # of required reducers.
2020-11-25 05:26:22,445 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Using reducer estimator: org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator
2020-11-25 05:26:22,459 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.InputSizeReducerEstimator - BytesPerReducer:1000000000 maxReducers=999 totalInputFileSize=107
2020-11-25 05:26:22,459 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting Parallelism to 1

```

```

2020-11-25 05:26:23,513 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - creating jar file Job340031370605982556.jar
2020-11-25 05:26:23,585 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - jar file Job340031370605982556.jar created
2020-11-25 05:26:23,406 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.JobControlCompiler - Setting up single store job
2020-11-25 05:26:23,467 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Key [pig.schematuple] is false, will not generate code.
2020-11-25 05:26:23,467 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Starting process to move generated code to distributed cache
2020-11-25 05:26:23,468 [main] INFO org.apache.pig.data.SchemaTupleFrontend - Setting key [pig.schematuple.classes] with classes to deserialize []
2020-11-25 05:26:23,449 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 1 map-reduce job(s) waiting for submission.
2020-11-25 05:26:23,454 [JobControl] INFO org.apache.hadoop.yarn.client.impl.Proxy - Connecting to ResourceManager at /6.0.0.8:8032
2020-11-25 05:26:23,477 [JobControl] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:26:23,758 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:26:23,758 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:26:23,762 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 05:26:23,769 [JobControl] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:26:23,769 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
2020-11-25 05:26:23,775 [JobControl] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths (combined) to process : 1
2020-11-25 05:26:23,849 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - number of splits:2
2020-11-25 05:26:23,940 [JobControl] INFO org.apache.hadoop.mapreduce.JobSubmitter - Submitting tokens for job: job_1606301819383_0021
2020-11-25 05:26:23,839 [JobControl] INFO org.apache.hadoop.yarn.client.impl.YarnClientImpl - Submitted application application_1606301819383_0021
2020-11-25 05:26:23,844 [JobControl] INFO org.apache.hadoop.mapreduce.Job - The url to track the job: http://quickstart.cloudera:8080/proxy/application_1606301819383_0021/
2020-11-25 05:26:23,844 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - HadoopJobId: job_1606301819383_0021
2020-11-25 05:26:23,844 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Processing aliases A,B,C,Z
2020-11-25 05:26:23,844 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - detailed locations: M: A[7,4],A[-1,-1],C[15,4],B[14,4],B[-1,-1],C[15,4] C: B: Z[-1,-1]
2020-11-25 05:26:23,844 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - More information at: http://localhost:50030/jobdet
ails.jsp?jobid=job_1606301819383_0021
2020-11-25 05:26:23,849 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 0% complete
2020-11-25 05:27:43,299 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 25% complete
2020-11-25 05:27:43,912 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 50% complete
2020-11-25 05:28:12,106 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - 100% complete
2020-11-25 05:28:12,106 [main] INFO org.apache.pig.tools.pigstats.SimplePigStats - Script Statistics:

HadoopVersion: 2.6.0-cdh5.12.0
PigVersion: 0.12.0-cdh5.12.0
User:Id: cloudera
FinishedAt: 2020-11-25 05:28:12
Features: 2020-11-25 05:28:12
COGROUP:

Success!

Job Stats (time in seconds):
JobId Maps Reduces MaxMapTime MinMapTime AvgMapTime MedianMapTime MaxReduceTime MinReduceTime AvgReduceTime MedianReduceTime Alias F
eature Outputs
job_1606301819383_0021 2 1 38 38 38 38 21 21 21 21 A,B,C,Z COGROUP hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/551033165
213/tmp551033165,
Input(s):
Successfully read 8 records from: "/user/cloudera/trainP/dummyf.txt"
Successfully read 8 records from: "/user/cloudera/trainP/dummyfb.txt"

Output(s):
Successfully stored 5 records (77 Bytes) in: "hdfs://quickstart.cloudera:8020/tmp/temp-1340503213/tmp551033165"

Counters:
Total records written : 5
Total bytes written : 77
Spillable Memory Manager spill count : 0
Total bags proactively spilled: 0
Total records proactively spilled: 0

Job DAG:
job_1606301819383_0021

2020-11-25 05:28:12,386 [main] INFO org.apache.pig.backend.hadoop.executionengine.mapReduceLayer.MapReduceLauncher - Success!
2020-11-25 05:28:12,389 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - fs.default.name is deprecated. Instead, use fs.defaultFS
2020-11-25 05:28:12,389 [main] INFO org.apache.hadoop.conf.Configuration.deprecation - mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address
2020-11-25 05:28:12,391 [main] INFO org.apache.pig.data.SchemaTupleBackend - Key [pig.schematuple] was not set... will not generate code.
2020-11-25 05:28:12,410 [main] INFO org.apache.hadoop.mapreduce.lib.input.FileInputFormat - Total input paths to process : 1
2020-11-25 05:28:12,411 [main] INFO org.apache.pig.backend.hadoop.executionengine.util.MapRedUtil - Total input paths to process : 1
{(1,(1,2),(1,2)))
(2,((2,4),(2,4)))
(4,((4,3)))
(5,((5,8)))
(23,((23,44),(23,44)))
grunts> █

```